

Optimal Control of a Wave Energy Converter via Recurrent Reinforcement Learning

HyunTaek Oh
School of EECS
Oregon State University
Corvallis, Oregon 97331
ohhyun@oregonstate.edu

Abstract—Despite the potential of wave energy as a sustainable resource, maximizing power capture in a Wave Energy Converter (WEC) remains a challenge, as conventional control strategies often fail to adapt to the inherent irregularity of real-world wave conditions. To address these limitations, this research proposes an AI-driven control architecture designed to maximize energy extraction by adaptively responding to highly irregular and dynamic wave conditions. The proposed architecture utilizes Proximal Policy Optimization (PPO) with an LSTM layer to achieve optimal energy extraction and real-time control efficiency. Experimental results on 200 irregular time-series wave datasets demonstrate that the proposed RL model achieves an average cumulative energy extraction of 3.38×10^8 , representing a 8.68% improvement over optimal fixed damping, and it is also comparable or better performance compared to the state-of-the-art controller. These findings suggest that the LSTM-PPO architecture provides a robust and superior alternative for real-time WEC optimization in dynamic marine environments.

Index Terms—Deep reinforcement learning; Energy harvesting; Long short-term memory (LSTM); Optimal control; Proximal policy optimization (PPO); Renewable energy; Wave energy converter (WEC).

I. INTRODUCTION

Ocean wave energy is highly viable renewable resources due to its superior predictability compared to solar and wind [1], [2], making it crucial for future global electricity demands. In the United States, the technical generation capacity of marine energy is estimated at a staggering 2,300 TWh/yr [3], with wave energy alone accounting for a dominant 1,400 TWh/yr, which is an amount equivalent to 34% of the total U.S. electricity generation. Wave Energy Converters (WECs) serve as the primary technology to harvest this dense resource. However, maximizing extraction efficiency remains a significant technical challenge due to the irregular, non-linear nature of real sea states with constantly varying frequencies and amplitudes [4]. To capture this energy effectively, the internal Power Take-Off (PTO) damping mechanism must be continuously and dynamically adjusted.

To address these environmental irregularities, traditional ocean engineering has widely adopted Model Predictive Control (MPC) [5]. MPC attempts to maximize energy extraction by forecasting future wave elevations, but its real-world performance is heavily constrained by its reliance on precise mathematical models of the WEC hydrodynamics. Formulating

these analytical models requires extensive domain expertise, and they are typically simplified to linear systems for real-time computational feasibility. Consequently, when subjected to the highly non-linear dynamics of irregular seas, these simplifications introduce modeling errors that degrade control performance and overall energy efficiency [6].

To overcome these limitations, model-free Reinforcement Learning (RL) has emerged as a strong alternative that bypasses complex mathematical modeling by interacting directly with the environment [7]. While early studies successfully applied deep RL to WECs [8], these previous models lacked the architectural capacity to handle sequential, time-series data. Because ocean waves are continuous and stochastic, an agent must interpret past wave patterns to make optimal decisions. To bridge this gap, this paper proposes a PPO-LSTM control architecture for the RM3 Wave Energy Converter. The LSTM structure processes sequential wave dynamics, while a custom carry-over system transfers observation and reward variables between training episodes to maintain system continuity. This approach allows the agent to dynamically maximize Energy Efficiency (EE) while strictly adhering to hard mechanical constraints.

The main contributions of this paper are as follows:

- We propose PPO-LSTM based control strategy that effectively processes the continuous and sequential dynamics of irregular ocean waves to dynamically adjust PTO damping.
- We introduce a custom carry-over logic that transfers observation and reward variables across split training episodes, ensuring continuous temporal learning without losing context.
- We design a novel EE-based reward function coupled with hard constraints to simultaneously maximize power extraction and guarantee mechanical safety.
- We provide the full implementation of our project as open-source to support future research (https://github.com/oh505/WEC_RL).

II. BACKGROUND

A. The RM3 Point Absorber

The Reference Model 3 (RM3) is a two-body wave energy converter consisting of a surface float and a submerged spar

[9]. As irregular ocean waves induce relative heave (vertical) motion between these two components, a Power Take-Off (PTO) system converts this kinetic energy into electricity. from a control systems perspective, the PTO functions as an adjustable damper. Dynamically tuning this damping coefficient serves as the agent's primary control action to maximize energy extraction from the environment.

However, operating WECs in dynamic oceanic conditions requires strict adherence to hardware safety margins. To prevent mechanical failure, the relative vertical displacement of the float must be rigidly constrained within a limit of $\pm 2.5m$ [9]. Consequently, the primary control objective is to maximize power extraction, subject to these strict mechanical boundaries [4]. Navigating the mechanical boundary within a stochastic wave environment motivates a specialized deep reinforcement learning architecture [8]. Specifically, Proximal Policy Optimization (PPO) is required for its stable policy updates in continuous action spaces, coupled with Long Short-Term Memory (LSTM) networks to effectively interpret time-series data and capture the continuous fluctuations of irregular waves.

B. Proximal Policy Optimization (PPO)

Applying deep reinforcement learning to real-time PTO damping control presents significant stability challenges, as traditional Policy Gradient (PG) methods are highly sensitive to step sizes [10]. In a stochastic wave environment, unconstrained parameter updates can cause drastic policy shifts and performance collapses, leading to critical violations of the WEC's strict physical displacement limits. To overcome these limitations, PPO adopts an Actor-Critic architecture [11], where the Actor network (π_θ) determines the continuous PTO damping coefficient and the Critic network estimates the expected state value ($V(s_t)$). The baseline expectation of the Critic is subtracted from the actual return via the Advantage function $A_t = R_t - V(s_t)$, ensuring that the agent evaluates the true effectiveness of its control actions independent of inherent environmental variance [12].

To mathematically restrict the step size of policy updates, PPO tracks the divergence between the current and previous policies using a probability ratio, $r_t(\theta)$, expressed as [13]:

$$r_t(\theta) = \frac{\pi_\theta(a_t|s_t)}{\pi_{\theta_{old}}(a_t|s_t)}$$

An unconstrained maximization of this ratio can lead to excessively large policy shifts. To ensure stability, PPO utilizes a clipped surrogate objective function [13]:

$$L^{CLIP}(\theta) = \mathbb{E}_t[\min(r_t(\theta)A_t, \text{clip}(\hat{r}_t(\theta), 1 - \epsilon, 1 + \epsilon)A_t)]$$

where ϵ is a hyperparameter defining the clipping threshold. The min operator establishes a conservative lower bound by penalizing the objective if the policy moves too far from the old one. This clipping mechanism forces the agent to optimize the PTO damping actions within a stable trust region, effectively preventing erratic control updates that could breach the physical safety boundaries of the system.

C. Long Short-Term Memory (LSTM) Networks

Long Short-Term Memory (LSTM) networks are designed to model long-term temporal dependencies while mitigating the vanishing gradient problem [14]. Unlike standard Multi-Layer Perceptrons (MLPs) that evaluate inputs instantaneously at a single time step and struggle to capture the continuous phase of irregular waves, LSTMs maintain an internal memory mechanism to preserve historical information across time steps [15]. Because ocean wave profiles exhibit highly non-linear and stochastic time-series characteristics, embedding an LSTM layer into our deep reinforcement learning framework is crucial. This structure enables the PPO agent to effectively process continuous sequences of historical observations, capturing underlying temporal contexts to maintain stable Power Take-Off (PTO) damping control under unpredictable wave conditions.

III. TECHNICAL APPROACH AND METHODS

A. System Overview: The RL-WEC Control Loop

The proposed system maximizes the energy extraction of a WEC in irregular ocean waves using Deep Reinforcement Learning (DRL). As shown in Figure 1, the architecture operates as a closed-loop control system. JONSWAP spectrum-based irregular waves drive the WEC simulator, which is strictly limited to a $\pm 2.5m$ displacement to ensure hardware safety. Before reaching the controller, an Observation Processing module filters the wave data and builds an observation history. To train the model on long time-series data without losing context, a custom carry-over logic transfers internal variables and observation histories between partitioned training episodes.

The main controller is a PPO-LSTM agent. The LSTM architecture helps the agent understand the time-dependent patterns of the waves. The agent outputs a damping control action, and its performance is scored by an Energy Efficiency reward function. This function compares the actual extracted power against the available wave energy flux, guiding the agent to find the most efficient and safe control strategy.




Fig. 1. Block diagram of the proposed PPO-LSTM control system for the Wave Energy Converter

B. Problem Formulation

To optimize the PTO damping control using a reinforcement learning approach, the interaction between the WEC and the irregular ocean environment is formulated as a Markov Decision Process (MDP). The environment simulates the WEC dynamics subjected to JONSWAP irregular waves, continuously providing states and evaluating actions based on physical boundaries.

Instead of depending on instantaneous values, the agent observes a state vector s_t that incorporates both current measurements and historical data to capture the wave's temporal dynamics. The observation space includes:

$$\begin{aligned} \text{state } (s_t) &= [z, \dot{z}, F_{exf}, F_{exp}, B_{pto,t-1}, \eta, \hat{\eta}] \\ \text{action } (a_t) &= B_{pto} \end{aligned}$$

Where $z, \dot{z}$ are the relative displacement and relative velocity between two bodies, F_{exf}, F_{exp} are the float and plate excitation forces, $B_{pto,t-1}$ is the previous damping coefficient prior to one step, and $\eta, \hat{\eta}$ are the current wave elevation and filtered current wave elevation, respectively. Inputs to the WEC plant are the float excitation force, plate excitation force, and PTO force.

To eliminate high-frequency noise and irrelevant high-bandwidth components, we introduce a causal band-pass filter. In real-time wave energy conversion, maintaining strict causality is essential because future wave data is unavailable during online operation [4]. This filtering process helps the reinforcement learning agent better capture the phase trends of the incoming waves. Standard reinforcement learning algorithms often suffer from training instability when exposed to raw, noisy wave observations [16]. We design a second-order band-pass filter using MATLAB's butter function, which results in a 4th-order transfer function. Therefore, the digital filter implementation maintains four past states for both inputs and outputs.

The low and high cut-off frequencies are calculated as follows:

$$\begin{aligned} \omega_{low} &= \omega_p * (1 - r) \\ \omega_{high} &= \omega_p * (1 + r) \end{aligned}$$

Where, $\omega_p = 0.644$ rad/s represents the peak frequency of Sea State 5, and r is the filter bandwidth ratio. The physical interpretation of this bandwidth ratio and its operation on the JONSWAP spectrum are visually demonstrated in Fig. 5 (Section V-C). Focusing the state space on the dominant wave excitation force zone ensures that the controller isolates the most energetic wave components [17]. The discrete-time filter is implemented using the Direct Form 1 difference equation:

$$y[n] = \frac{1}{a_1} \left(\sum_{i=0}^4 b_{i+1} x[n-i] - \sum_{j=1}^4 a_{j+1} y[n-j] \right)$$

The action a_t chosen by the agent corresponds to the PTO damping coefficient (B_{pto}). To maintain operational stability, the continuous action space is bounded within a range of $[0 \text{ to } 1.2 \times 10^7] N \cdot s/m$.

To enable the sequence-learning capability of the LSTM, a history window of 100 time steps is stacked. Furthermore, to maintain context across segmented training episodes, the internal variables of the reward function and the observation history are carried over to subsequent segments. This ensures continuous sequential awareness without breaking the physical timeline.

A strict physical displacement constraint is imposed on the WEC floats since they cannot oscillate indefinitely. The vertical displacement (z) is bounded to exactly $\pm 2.5m$. If the agent attempts a control strategy that drives the float beyond this $\pm 2.5m$ limit, the action is penalized, preventing the policy from exploiting unrealistic mechanical behaviors.

C. Physics-based Reward Formulation

The objective of the WEC control system is to maximize energy extraction. However, defining the reward simply as the raw extracted energy is suboptimal in an irregular wave environment, as the absolute energy potential fluctuates depending on the varying scales of the incoming waves. To normalize the learning signal across diverse wave states, a physics-based Energy Efficiency (EE) reward function is proposed.

Instead of accumulating the extracted energy, the reward formulation focuses on the comparison of average power. The Energy Efficiency is defined as the ratio of the time-averaged extracted power from the PTO system to the available wave energy flux over a given evaluation window:

$$R_{EE} = \frac{\bar{P}_{pto}}{\bar{P}_{wave}} = \frac{\text{Average Extracted Power}}{\text{Available Wave Energy Flux}}$$

This normalization ensures the PPO-LSTM agent learns a generalized optimal control policy that maximizes extraction efficiency regardless of the wave magnitude.

In addition, to enforce the physical safety of the WEC, the final reward function R_t integrates a penalty mechanism. If the agent's action causes the relative displacement to exceed the strict $\pm 2.5m$ mechanical limit, a negative reward is applied:

$$R_t = \begin{cases} R_{EE}, & \text{if } |z| \leq 2.5 \\ \text{Penalty}, & \text{if } |z| > 2.5 \end{cases}$$

By combining the normalized power efficiency with the boundary penalty, the reward function prevents the agent from finding reward hacking strategies, which guides it toward a physically viable and efficient control policy.

D. PPO-LSTM Agent Architecture

To achieve continuous and stable control over the PTO damping coefficient, the Proximal Policy Optimization (PPO) algorithm is used. While standard Multi-Layer Perceptron (MLP) networks evaluate instantaneous observations, the irregular dynamics of JONSWAP waves require temporal awareness. To overcome this, a unified PPO-LSTM architecture is designed, as illustrated in Figure 2. With this structure, the agent can capture temporal dependencies and optimize energy extraction. The data processing flows through the network in three progressive stages. First, the shared input




Fig. 2. PPO-LSTM Actor-Critic Architecture

layer, consisting of the 7 observation features (s_t), is fed independently into the Actor and Critic networks. In both branches, the raw observation features are processed through fully-connected (FC) layers with ReLU activation to extract high-dimensional, temporally invariant feature representations. Next, the extracted feature sequences are passed into the LSTM layer (256 units). The recurrent structure of the LSTM processes the sequences sequentially, capturing the temporal context of the wave dynamics and the corresponding WEC states. This ensures the agent understands how the environment gradually changes over the given history window. Finally, the outputs of the LSTM are mapped to their respective final values via the back-end FC layers. In the Actor network, the architecture branches into two paths to compute the mean (μ) and standard deviation (σ) of the policy distribution. The mean branch utilizes a Tanh function with a custom Scaling layer, while the standard deviation branch employs a Softplus function. These parameters collectively determine the final sampled continuous damping action ($a_t = B_{pto}$). Concurrently, the Critic network utilizes a single linear FC layer to estimate the state value (V_t). It evaluates the overall quality of the observation sequence to guide the Actor's policy updates.

IV. EXPERIMENTAL SETUP

This section details the simulation environment, the wave dataset configuration, the training strategies, and the hyperparameter settings used to evaluate the proposed control framework.

A. Simulation Platform and WEC Model

The entire closed-loop control system, including the WEC hydrodynamics, mechanical constraints, and the DRL agent, is implemented within the MATLAB and Simulink (R2025a) environment. The WEC plant is modeled after Reference Model 3 (RM3), a standard two-body floating point absorber. To achieve high-fidelity control and capture precise physical

interactions, the simulation operates at a sampling frequency of 10Hz , corresponding to a fixed time step of $\Delta t = 0.1\text{s}$. To manage the heavy computational demands of simulating complex marine environments over long time horizons, training is accelerated via High-Performance Computing (HPC) cluster managed by a Slurm scheduler, utilizing parallel workers across dedicated CPU and GPU resources.

B. Wave Environment Configuration (Sea State 5)

To test the PPO-LSTM agent in realistic conditions, an irregular wave dataset is generated using the JONSWAP spectrum. The simulation focuses on Sea State 5, which is a common operating environment for wave energy converters. The specific wave parameters are listed in Table I. The dataset contains 1,000 unique wave time-series, each lasting 1,200 seconds. The data is split into 80% for training and 20% for validation.

TABLE I
PARAMETERS FOR JONSWAP WAVE SPECTRUM (SEA STATE 5)

Parameter	Value	Unit
Significant Wave Height (H_s)	3.66	m
Peak Period (T_p)	9.7	s
Rho (ρ)	1000	kg/m^3
Gravitational Acceleration (g)	9.81	m/s^2
Peak Enhancement Factor (γ)	3.3	-
Simulation Duration	1,200	s
Time Step (Δt)	0.1	s

C. Training Strategy and Carry-Over Logic

At a 10Hz sampling rate, a full 1,200-second wave sequence generates 12,000 time steps. Training on this entire sequence at once causes unstable gradient updates. To optimize training Efficiency, each wave is partitioned into shorter sub-episodes of 200 steps (20 seconds). This length was empirically selected to balance gradient stability and the effective learning horizon determined by the discount factor. The agent was trained for a total of 48,000 sub-episodes.

To prevent physical discontinuities at the segment boundaries, a variable carry-over mechanism is introduced. Instead of transferring complex LSTM hidden states, the environment resets them but passes reward and observation variables to the next sub-episodes. Specifically, the final step of the state observation history and the exact zero-crossing detection states are saved and used to initialize the subsequent segment. This mechanism ensures the agent perceives the partitioned episodes as a single, continuous physical timeline, which is essential for accurate reward assignment and stable learning in irregular waves.

D. Hyperparameter Configuration

To guarantee the algorithmic validity and reproducibility of the trained control policy, the hyperparameters configured within the MATLAB Reinforcement Learning Toolbox are

defined. These parameters dictate the optimization landscape and convergence properties of the PPO-LSTM architecture, as summarized in Table II.

TABLE II
HYPERPARAMETERS FOR PPO-LSTM TRAINING

Parameter	Value
Experience Horizon	400
Learning Frequency	400
Mini Batch Size	100
Discount Factor	0.99
Number of Epochs	5
Entropy Loss Weight	0.001
Actor Learning Rate	1×10^{-3}
Actor Gradient Threshold	5
Critic Learning Rate	1×10^{-3}
Critic Gradient Threshold	5
Sample Time	0.1

V. RESULTS

A. Benchmark Baselines and Evaluation Metrics

To evaluate the performance of the proposed PPO-LSTM controller, two distinct baseline control strategies from traditional ocean engineering are established for comparative analysis. For a fair evaluation, all simulations, including the baselines and the proposed learning-based agent, are executed with a time-step of 0.1s. Furthermore, all control frameworks are restricted to damping-only control, ensuring that performance gains are driven by behavioral control logic rather than hardware or stiffness variations.

- **Fixed-Damping Control (Baseline 1):** The conventional passive control strategy where the damping coefficient (B_{pto}) remains constant throughout the test duration. As illustrated in Figure 3, an exhaustive brute-force grid search was conducted over the range of 0 to $12MN \cdot s/m$ with a step size of $0.3MN \cdot s/m$. The optimal damping coefficient was determined to be $5.4MN \cdot s/m$, yielding a maximum mean cumulative energy extraction of $3.11 \times 10^8 J$, which serves as the primary benchmark.
- **Instantaneous Frequency (IF) Control (Baseline 2):** A state-of-the-art non-learning active control method that adaptively adjusts the damping command based on the real-time wave frequency. For a fair comparison, we disabled the stiffness control in the IF algorithm so that it relies on damping control, matching the conditions of the other models.

The control efficiency and power extraction performance are quantified using the following two metrics, evaluated over a validation set consisting of 200 distinct irregular wave scenarios:




Fig. 3. Mean cumulative absorbed energy across 200 validation wave scenarios as a function of the PTO damping coefficient (B_{pto}). The optimal baseline performance is achieved at $B_{pto} = 5.4 MN \cdot s/m$, yielding a maximum energy of $3.11 \times 10^8 J$.

- 1) **Mean Cumulative Absorbed Energy:** The average of the total electrical energy extracted at the end of each evaluation episode across the 200 test wave profiles.
- 2) **Capture Width Ratio (CWR):** A non-dimensional efficiency metric commonly used in wave energy conversion to evaluate how effectively the device captures energy relative to the incoming wave power. It is defined as:

$$CWR = \frac{P_{absorbed}}{P_{incident} \cdot B}$$

where $P_{absorbed}$ represents the average power captured by the WEC, $P_{incident}$ is the wave energy flux per unit crest length, and $B = 20$ denotes the width of the device.

B. Training Convergence and Reward Function Analysis

This section analyzes how the reward function design affects training stability. For this analysis, a basic Multi-Layer Perceptron (MLP) architecture is used. By adopting a simplified Multi-Layer Perceptron (MLP) architecture without recurrent elements, we intentionally exclude the structural advantages of the LSTM layer, such as its inherent temporal noise-filtering capabilities. This approach allows us to isolate and independently validate the mathematical convergence and stability of the proposed physics-based EE reward function itself. The observation space is simplified to include relative displacement (z), relative velocity ($\dot{z}$), wave excitation forces (F_{exf} , F_{exp}), previous damping value, and current wave elevation.

Although the model was trained for a total of 48,000 episodes, Fig. 4 plots only the first 10,000 episodes. The training curves stabilized completely after this point, showing no significant changes in the remaining episodes.

In Fig. 4, the proposed physics-based EE reward (blue) is compared against an alternative linear reward function (red). Please note that this linear reward is simply a different reward design used for comparison, not the fixed-damping control




Fig. 4. Training convergence curves for the MLP architecture over the first 10,000 episodes. The proposed EE reward (blue) shows much better stability and lower variance than the alternative linear reward (red).

baseline described in Section 5.A. Specifically, the linear reward is calculated by scaling the raw extracted power by 3×10^{-6} and subtracting the same boundary penalty term used in the EE reward.

The linear reward shows a very wide reward interval. This is because the linear formulation relies on raw extracted power, which changes wildly depending on the size of the waves. In contrast, the proposed EE reward maintains a very tight and consistent interval. By dividing the captured power by the incoming wave energy flux, the EE reward removes the effect of wave scales and gives the agent steady feedback.

The sharp drop in the blue curve before episode 500 shows the agent quickly learning to avoid basic physical penalties and boundaries. After this initial phase, the EE reward stays within a stable range. The continuous minor fluctuations in the curve are caused by the random nature of the JONSWAP wave spectrum. Because each episode mixes different wave sizes randomly, some episodes are naturally harder to control than others. This leads to small ups and downs in the learning curve.

C. Ablation Study: Impact of Wave Elevation Filtering

To validate the necessity of the proposed causal band-pass filter, an ablation study was conducted. Since the Long Short-Term Memory (LSTM) architecture inherently possesses the ability to filter temporal noise through its internal gating mechanisms, evaluating the signal filter within the PPO-LSTM framework could obscure the filter's independent contribution. Therefore, to isolate the pure effect of data processing, this study utilized a standard Multi-Layer Perceptron (MLP) agent without memory. The observation space was kept identical, and the agent was trained across five independent random seeds under different filter bandwidth ratios ($r \in \{0.1, 0.2, 0.25, 0.5\}$) alongside a 'No Filter' baseline.

Before evaluating the control performance across these ratios, Fig. 5 shows the practical meaning of the bandwidth ratio (r) in the frequency domain. In the energy spectrum of Sea State 5, most of the wave's power is concentrated around




Fig. 5. Energy density spectrum of Sea State 5 comparing raw and filtered signals ($r = 0.2$).

the peak frequency ($\omega_p = 0.644$ rad/s). By setting a specific filter range, such as the $r = 0.2$ area highlighted in the figure, we can successfully capture this main chunk of energy. This filter intentionally blocks out useless, slow moving drift and fast, choppy noise. As a result, the RL agent receives a much cleaner observation state to learn from. However, finding the right width for this filter is a balancing act. It must be narrow enough to remove confusing noise, but wide enough to keep the important shape and timing (phase) of the actual waves.




Fig. 6. Distribution of mean extracted energy across 200 validation wave scenarios for the 'No filter' baseline and different filter bandwidth ratios (r).

As illustrated in Fig. 6 and Fig. 7, the 'No Filter' baseline fails to maintain physical constraints, yielding a median survival rate below 30%. This confirms that unfiltered high-frequency noise directly induces erratic control actions and boundary violations. To identify the optimal filter configuration, various bandwidth ratios (r) were evaluated. A narrow band ($r = 0.1$) over-smooths the signal, losing critical phase information and crashing the survival rate. Conversely, wider bands ($r \geq 0.25$) allow excessive noise to degrade policy stability, resulting in high variance. Only $r = 0.2$ maintains the optimal balance, consistently achieving a near 100% survival




Fig. 7. Survival rate representing the compliance with the $\pm 2.5m$ physical displacement constraint.

rate across all random seeds while maximizing the median extracted energy. Consequently, the causal band-pass filter with $r = 0.2$ was selected as the optimal pre-processing module and integrated into the primary PPO-LSTM architecture. The performance of this fully optimized model is evaluated in the following sections.

Interestingly, while the MLP without memory mechanism requires a strict bandwidth constraint ($r = 0.2$) to maintain safety, the primary PPO-LSTM architecture optimizes its performance at a wider bandwidth ($r = 0.25$). This indicates that the recurrent gating mechanism of the LSTM effectively dampens residual high-frequency irregularities, allowing the controller to utilize richer sequential features for higher energy extraction without sacrificing mechanical safety. Consequently, $r = 0.25$ was established as the filter configuration for the main architecture evaluated in the subsequent benchmark analyses.

D. Control Performance Comparison

This section analyzes the power extraction efficiency of the proposed PPO-LSTM controller compared to the benchmark strategies and the theoretical upper limit across 200 irregular wave scenarios. The exact values of the mean cumulative absorbed energy and the corresponding percentage improvements over the Fixed-Damping baseline are summarized in Table III.

TABLE III
PERFORMANCE SUMMARY OF CONTROL METHODS

Control Method	Mean Energy (10^8 J)	Improvement (%)
Theoretical Max	6.30	-
Fixed-Damping (Baseline)	3.11	Base
Instantaneous Frequency (IF)	3.37	+8.36%
Proposed PPO-LSTM	3.38	+8.68%

As shown in Table III, both active control methods extract more energy than the passive Fixed-Damping baseline (3.11×10^8 J). Because the passive baseline relies on a single constant

damping coefficient, its performance is limited in fluctuating, stochastic wave environments, which underscores the necessity of active control.

The best-performing PPO-LSTM agent, selected from five independent training runs, achieves an average energy extraction of 3.38×10^8 J (an 8.68% increase over the baseline), which is comparable to the state-of-the-art IF active control method (3.37×10^8 J, 8.36% improvement). This result indicates that the learning-based framework can closely match the performance of advanced ocean engineering control strategies. Notably, the PPO-LSTM agent achieves this near state-of-the-art efficiency without requiring real-time wave frequency calculations during operation, presenting a distinct practical advantage.

To further evaluate the statistical robustness of the controllers across all 200 test scenarios, the distribution of the Capture Width Ratio (CWR) is presented in Fig. 8.




Fig. 8. Comparison of Capture Width Ratio (CWR) distributions across 200 validation irregular wave scenarios.

As illustrated in Fig. 8, the CWR distribution of the PPO-LSTM controller closely aligns with that of the advanced IF method. While there are no prominent statistical deviations between the two active control distributions, the overall efficiency profiles of both methods show an upward shift compared to the passive baseline.

The fact that the learning-based model replicates the statistical efficiency distribution of a state-of-the-art engineering controller is highly meaningful. It demonstrates that even with a straightforward neural network structure, the agent effectively captures the sequential features of the irregular waves. The proposed model provides a reliable and practical alternative that yields robust efficiency across diverse JONSWAP wave profiles.

E. Dynamic Control Behavior Analysis

To understand how the controllers operate in real-time, a time-series analysis is conducted using a representative irregular wave segment (Wave 149). Fig. 9 shows the wave elevation, the damping command of the proposed PPO-LSTM model, and the damping command of the IF method over the same time period.




Fig. 9. Time-series control behavior under irregular wave scenario 149: (a) wave elevation, (b) damping coefficient command of the IF controller, and (c) damping coefficient command of the proposed PPO-LSTM controller.

As shown in Fig. 9(c), the proposed PPO-LSTM controller adjusts the damping coefficient smoothly and continuously. It keeps the control values low (under $1.1 \times 10^7 \text{ N} \cdot \text{s/m}$) and avoids sudden jumps. Because the LSTM layer remembers past data, the network can follow wave trends and make stable, proactive adjustments.

In contrast, Fig. 9(b) shows that the IF method produces sharp and sudden spikes. To react to instant frequency changes, it applies rapid impulses that reach much higher values, up to $5 \times 10^7 \text{ N} \cdot \text{s/m}$.

This comparison shows a major practical advantage of the proposed model. The PPO-LSTM agent achieves similar energy efficiency to the IF method but maintains a much smoother and lower control profile. This significantly reduces physical stress on the mechanical parts of the device.

VI. SUMMARY AND FUTURE WORK

This paper presents a reinforcement learning-based control strategy to increase the energy harvesting efficiency of wave energy converters (WECs). The proposed agent learns to adjust the PTO damping coefficient at the exact timing under irregular wave conditions. Using a straightforward neural network structure, the model achieved a mean cumulative energy extraction of $3.38 \times 10^8 \text{ J}$ and a Capture Width Ratio (CWR) of 27.1%. These results are comparable to state-of-the-art active control methods in ocean engineering. This highlights the promising capability of learning-based frameworks to provide adaptive and efficient control without requiring online frequency calculations.

Future research will investigate alternative neural network architectures, such as Transformers, as well as other reinforcement learning algorithms to further improve control performance and computational efficiency. The proposed control policy will also be evaluated under more challenging environmental conditions, including mixed wave profiles and

extreme sea states, to assess its robustness and reliability. In addition, a hybrid control framework that integrates the existing instantaneous frequency control strategy with the reinforcement learning policy will be developed to further enhance power capture performance.

ACKNOWLEDGMENT

This work was conducted as a joint project at Oregon State University. The author gratefully acknowledges Dr. Alan Fern and William Solow (School of Electrical Engineering and Computer Science) for their support on the AI models, Professor Solomon Yim (School of Civil and Construction Engineering) and Dr. In-Yong Kim (School of Electrical Engineering and Computer Science) for providing the WEC simulation model and valuable insights into the wave energy conversion system.

REFERENCES

- [1] A. F. d. O. Falcão, "Wave energy utilization: A review of the technologies," *Renewable and Sustainable Energy Reviews*, vol. 14, pp. 899–918, Apr. 2010.
- [2] G. Reikard, B. Robertson, and J.-R. Bidlot, "Combining wave energy with wind and solar: Short-term forecasting," *Renewable Energy*, vol. 81, pp. 442–456, Sept. 2015.
- [3] L. Kilcher, M. Fogarty, and M. Lawson, "Marine Energy in the United States: An Overview of Opportunities," 2021.
- [4] J. V. Ringwood, G. Bacelli, and F. Fusco, "Energy-Maximizing Control of Wave-Energy Converters: The Development of Control System Technology to Optimize Their Operation," *IEEE Control Systems Magazine*, vol. 34, pp. 30–55, Oct. 2014.
- [5] N. Faedo, S. Olaya, and J. V. Ringwood, "Optimal control, MPC and MPC-like algorithms for wave energy systems: An overview," *IFAC Journal of Systems and Control*, vol. 1, pp. 37–56, Sept. 2017.
- [6] E. Anderlini, D. I. M. Forehand, E. Bannon, Q. Xiao, and M. Abusara, "Reactive control of a two-body point absorber using reinforcement learning," *Ocean Engineering*, vol. 148, pp. 650–658, Jan. 2018.
- [7] E. Anderlini, "Control of wave energy converters using machine learning strategies," 2017.
- [8] E. Anderlini, S. Husain, G. G. Parker, M. Abusara, and G. Thomas, "Towards Real-Time Reinforcement Learning Control of a Wave Energy Converter," *Journal of Marine Science and Engineering*, vol. 8, p. 845, Nov. 2020.
- [9] V. Neary, Y.-H. Yu, J. Epler, K. Shoel, M. Previsic, M. Lawson, and Y. Li, "Marine and Hydrokinetic Data Repository (MHKDR)," Sept. 2014.
- [10] P.-W. Chou, D. Maturana, and S. Scherer, "Improving Stochastic Policy Gradients in Continuous Control with Deep Reinforcement Learning using the Beta Distribution," in *Proceedings of the 34th International Conference on Machine Learning*, pp. 834–843, PMLR, July 2017.
- [11] V. Konda and J. Tsitsiklis, "Actor-Critic Algorithms," in *Advances in Neural Information Processing Systems*, vol. 12, MIT Press, 1999.
- [12] J. Schulman, P. Moritz, S. Levine, M. Jordan, and P. Abbeel, "High-Dimensional Continuous Control Using Generalized Advantage Estimation," June 2015.
- [13] J. Schulman, F. Wolski, P. Dhariwal, A. Radford, and O. Klimov, "Proximal Policy Optimization Algorithms," July 2017.
- [14] S. Hochreiter and J. Schmidhuber, "Long Short-Term Memory," *Neural Computation*, vol. 9, pp. 1735–1780, Jan. 1997.
- [15] E. Anderlini, D. I. M. Forehand, E. Bannon, and M. Abusara, "Reactive control of a wave energy converter using artificial neural networks," *International Journal of Marine Energy*, vol. 19, pp. 207–220, Sept. 2017.
- [16] S. Zou, X. Zhou, I. Khan, W. W. Weaver, and S. Rahman, "Optimization of the electricity generation of a wave energy converter using deep reinforcement learning," *Ocean Engineering*, vol. 244, p. 110363, Jan. 2022.
- [17] D. García-Violini, N. Faedo, F. Jaramillo-Lopez, and J. V. Ringwood, "Simple Controllers for Wave Energy Devices Compared," *Journal of Marine Science and Engineering*, vol. 8, p. 793, Oct. 2020.